

nstructions

The following is a pure Python implementation of a very cool probabilistic data structure called HyperLogLog:

<https://github.com/svpcom/hyperloglog>

(Some research paper are linked at the bottom of the webpage in case you are interested in understanding how it works.)

Your task for this HW is to optimize the `get_nearest_neighbors` function in the file `hll.py`.

As usual, you can do this HW in teams of size up to 3.

Please submit a single modified `hll.py` file that implements the changes.